

CSCI 1515: Applied Cryptography

P. Miao

Spring 2026

These are lecture notes for CSCI 1515: Applied Cryptography taught at BROWN UNIVERSITY by Peihan Miao in the Spring of 2026.

These notes were originally taken by Jiahua Chen in Spring 2023, and were updated in Spring 2024 by Sudatta Hor. They are currently being maintained by Chloe Qiao and Ian Hajra. There has been gracious help and input from classmates and fellow TAs. Please direct any mistakes/errata to a thread on Ed, or feel free to pull request or submit an issue to the [notes repository](#).

Notes last updated January 30, 2026.

Contents

1	January 21, 2026	3
1.1	Introduction	3
1.1.1	Staff	3
1.1.2	Course Philosophy and Logistics	3
1.2	What is cryptography?	4
1.3	Secure Communication	5
1.3.1	Message Secrecy	6
1.3.2	Message Integrity	8
1.3.3	Projects	10
1.4	Zero-Knowledge Proofs	10
1.5	Secure Multi-Party Computation	12
1.6	Fully Homomorphic Encryption	15
1.7	Further Topics	16
1.8	Q & A	17
2	January 28, 2026	18
2.1	Encryption Scheme Basics	18
2.1.1	Syntax	19
2.1.2	Symmetric-Key Encryption Schemes	20
2.1.3	Public-Key Encryption Schemes	26
2.1.4	RSA	29

3 January 30, 2026	31
3.1 Basic Number Theory, <i>continued</i>	31
3.2 RSA Encryption, <i>continued</i>	31
3.3 Intro to Group Theory	33
3.4 Computational Assumptions	34
3.5 ElGamal Encryption	35
3.6 Secure Key Exchange	36

§1 January 21, 2026

§1.1 Introduction

The course homepage is at <https://cs.brown.edu/courses/csci1515/spring-2026/>, where you can find information such as the [syllabus](#), projects, homeworks, calendar, lectures and more.

The course is offered in-person in *Friedman 108*, as well as synchronously over Zoom and recorded asynchronously (lectures posted online). Lecture attendance and participation is highly encouraged!

EdStem will be used for course questions, and **Gradescope** is used for assignments.

§1.1.1 Staff

Peihan has been at Brown for a couple of years and this is the third time she is teaching this course. Before Brown, she was at the University of Illinois Chicago. Before that, she finished her PhD at UC Berkeley in 2019 with a focus in cryptography. Afterwards, she worked in industry for a couple of years (Visa) before deciding to come back to academia. She still collaborates with industry to see what problems need to be solved in practice.

During her PhD, she started off doing more theoretical cryptography but also did internships and found cryptography fascinating as well. Now she works in both.

Our course staff all have strong backgrounds in cryptography and are excited to help you learn!

§1.1.2 Course Philosophy and Logistics

If you look up other *applied* cryptography courses online or at other universities, you will find courses that have “applied” in their title. However, if you look at their syllabus or content, it’s still mostly theoretical crypto. This may (1) deter students from learning about crypto and (2) leave a gap between theoretical crypto and crypto in practice. (2) is bad because if someone makes a mistake in the crypto domain, the consequences are often significant.

As such, it’s helpful for students to get hands-on experience with cryptography:

- How cryptography has been used in practice and
- How cryptography will be used and implemented in the future.

The closest similar course is found at Stanford, which covers theoretical crypto in the first half and more applied crypto in the second. But even that course only covers very basic crypto that are very

well established. In the past 10 years or so, there are new and exciting topics in crypto that are gradually becoming more and more common which we will also cover in this course.

For this course, it will be *much less* about math and proofs, and much more about how you can use these tools to do something more fun. It will be coding heavy and all projects will be implemented in C++ using crypto libraries.

If, however, you are interested in the theoretical or mathematical side, you might consider other courses at Brown like CSCI 1510 and MATH 1580.

Capstone requirements are cancelled in FY25–26 for CS and joint-CS concentrators.

The following is the grading policy:

Type	Percentage
Introductions	1%
Project 0	4%
Projects 1 & 2 & 5	24% (8% each)
Projects 3 & 4	24% (12% each)
Code Review	8% (4% each)
Homeworks	25% (5% each)
Final Project	14%

You have a total of four (4) late days that you can use to extend the due dates of *Projects 0–5* free of charge, but at most two (2) late days may be applied to any one project. Beyond that, you will be penalized 40% of the project's value for each day it is late. There are **no** late days for *homeworks*, as we want to release the solution guide as soon as possible.

All projects are independent, but collaboration is allowed and encouraged. However, you *must write up your own code*.

If you're sick, let Peihan know with a Dean's note. For a full description of the late policy, refer to the syllabus.

§1.2 What is cryptography?

At a high level, *cryptography is a set of techniques that protect sensitive or important information*.

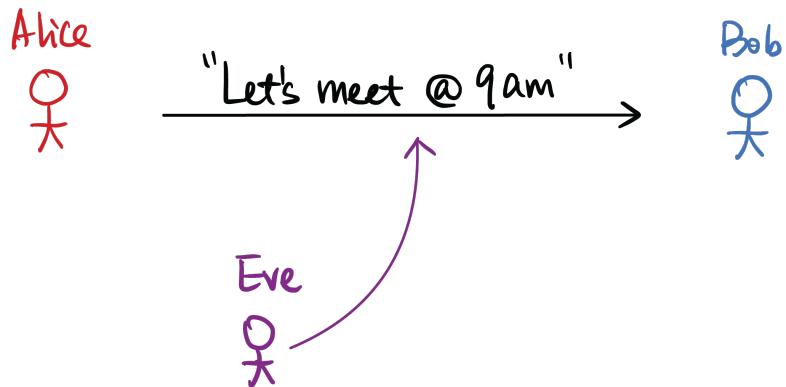
Question. Where is cryptography used in practice? What guarantees do we want in these scenarios?

- Online transactions

- When you make a purchase, you might not want people to see your bank balance, what else you have purchased, etc.
- You also want to ensure that it was really *you* who purchased the item and not somebody else i.e. authentication
- Secure messaging
 - End-to-end texting, iMessage
 - We don't want anyone else to see our messages
- Online voting
 - Privacy of votes, validity of votes
- Databases
 - Secure storage
- There is much more!

§1.3 Secure Communication

We'll start with the most classic form of cryptography: *secure communication*.



Assume Alice wants to communicate to Bob “Let's meet at 9am”, what are some security guarantees we want?

- Eve cannot *see* the message from Alice to Bob.

- Eve cannot *alter* the message from Alice to Bob.

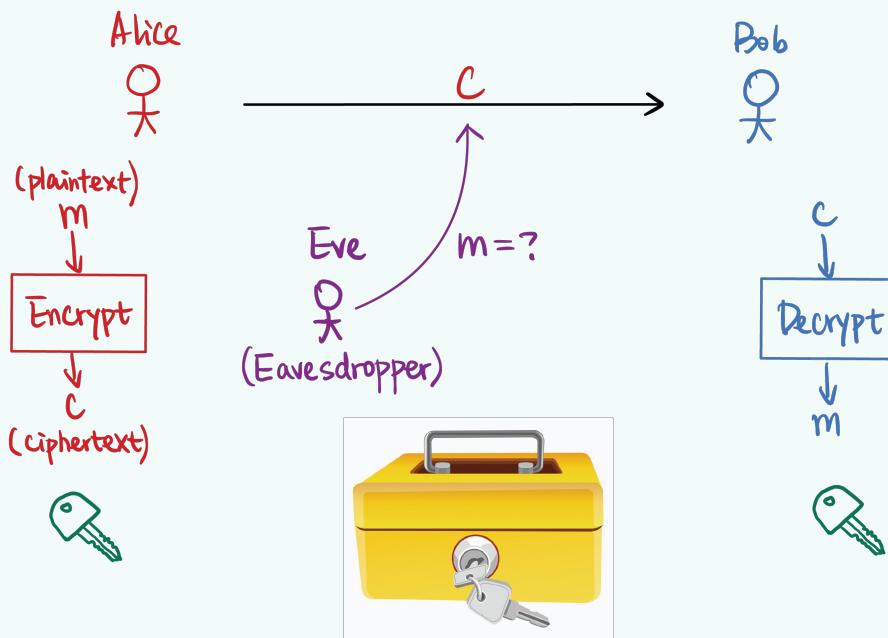
These two guarantees are the most important guarantees! The former is called message secrecy, the latter is called message integrity.

§1.3.1 Message Secrecy

Definition 1.1 (Message Secrecy)

We want cryptography to allow Alice to *encrypt* the message m (which we call *plaintext*) by running an algorithm that produces a *ciphertext* c . We call this an *encryption scheme*.

Bob will be able to receive the ciphertext c and run a *decrypt* algorithm to produce the message m again. This is akin to a secure box that Alice locks up, and Bob unlocks, while Eve does not know the message. The easiest way is for Alice and Bob to agree on a shared secret key.

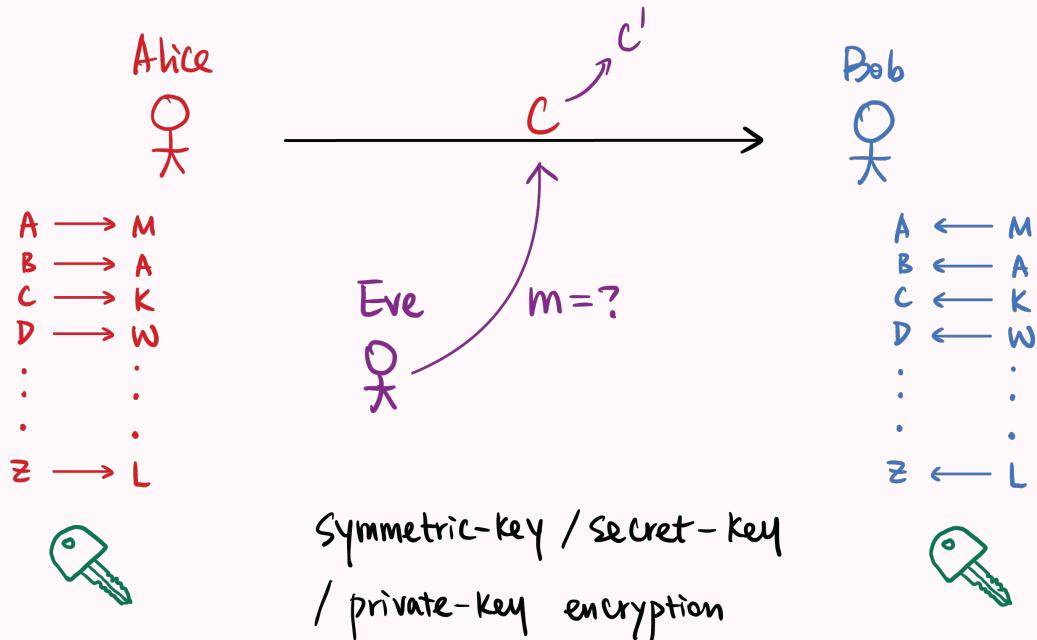


In this model, Eve is a weaker adversary, an *eavesdropper*. Eve can only see the message, not alter it.

Example 1.2 (Substitution Cipher)

The key that Alice and Bob jointly uses is a permutation mapping from $\{A \dots Z\} \rightarrow \{A \dots Z\}$. This mapping is the *secret key*.

Bob also has the mapping, and takes the inverse of the permutation to retrieve the message.



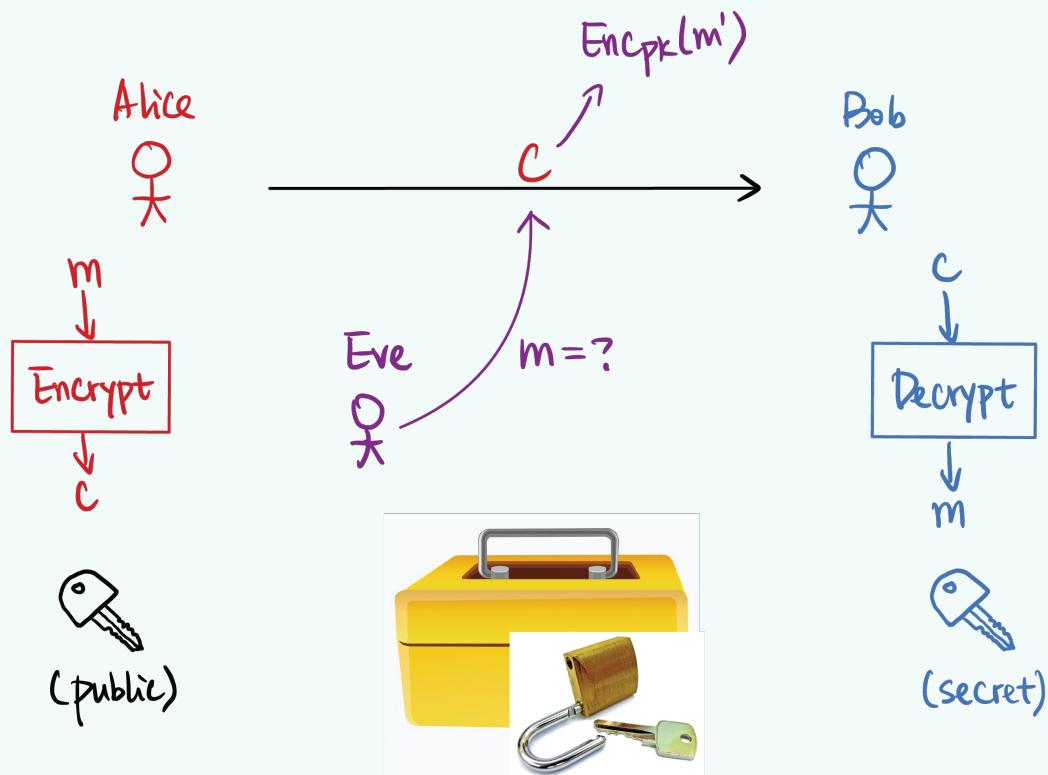
This scheme is not quite secure! *Why?*

Given a large enough text, you can apply frequency analysis and break the substitution cipher.

Remark. This encryption scheme also requires that Alice and Bob meet up in person to exchange this shared private key. Schemes like this are called *symmetric-key*, *secret-key*, or *private-key encryption*. They need to somehow agree first on the same secret key.

Definition 1.3 (Public-key Encryption)

There is another primitive that is much stronger: public-key encryption. Bob generates both a *public* key and a *private* key, and then publishes his public key. You can consider a lock where you don't need a key to lock it¹, and only Bob has the key to unlock it.

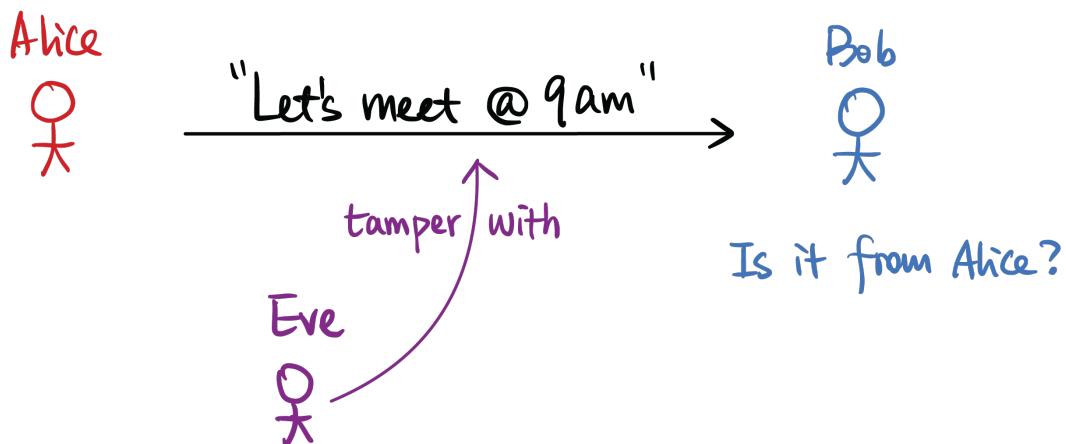


This is seemingly magic! Bob could publish a public key on his homepage, anyone can encrypt using a public key but only Bob can decrypt. *Stay tuned, we will see public-key encryption schemes next lecture!*

§1.3.2 Message Integrity

Alice wants to send a message to Bob again, but Eve is stronger! Eve can now tamper with the message.

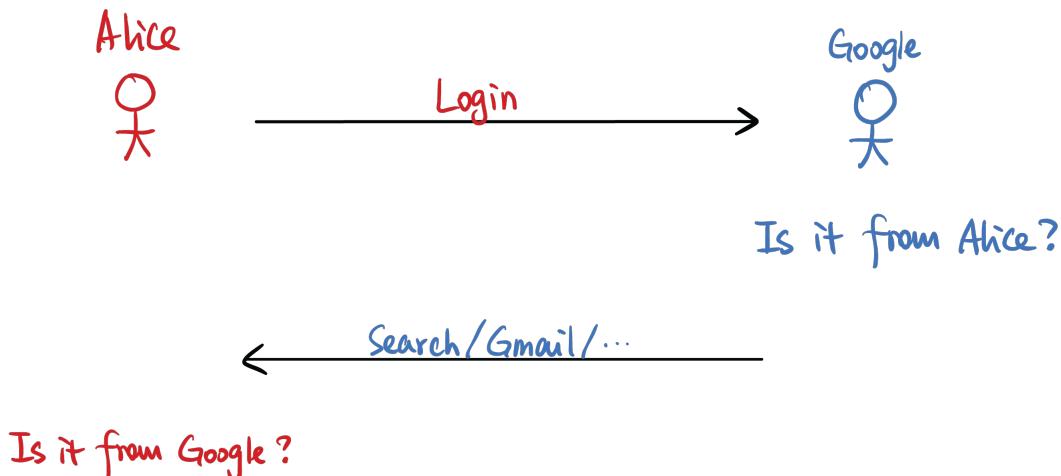
¹You literally click it closed



Bob wants to ensure that the message *actually* comes from Alice. Does our previous scheme (of encrypting messages) solve this problem? Nope!

Eve can change the ciphertext to something else, they could pretend to be Alice. In secret-key schemes, if Eve figures out the secret-key, they can forge messages from Alice. Even if Eve doesn't know the underlying message, they could still change it to some other ciphertext which might be correlated to the original ciphertext, *without knowing the underlying message*. We'll see how Eve can meaningfully do this in some schemes. Alice could send a message "Let's meet at x AM" and Eve could tamper this to say "Let's meet at $x + 1$ AM."

This is sort of an orthogonal problem to message secrecy. For example, when Alice logs in to Google, Google needs to verify that Alice actually is who she claims to be.



This property that we want is called message integrity.

§1.3.3 Projects

Projects Overview

0. *Cipher* Warm-up, you will implement some basic cryptographic schemes.
1. *Signal* Secure Communication: how to communicate in secret.
2. *Auth* Secure Authentication: how to authenticate yourself.
3. *Vote* Zero-Knowledge Proofs: we'll use ZKPs to implement a secure voting scheme.
4. *Yaos* Secure Multiparty Computation: we'll implement a way to run any function securely between two parties.
5. *PIR* Private Information Retrieval (using homomorphic encryption, a form of post-quantum cryptography).

We'll now introduce the latter three projects!

§1.4 Zero-Knowledge Proofs

This is to prove something without *revealing* any additional knowledge.

For example, Alice may want to

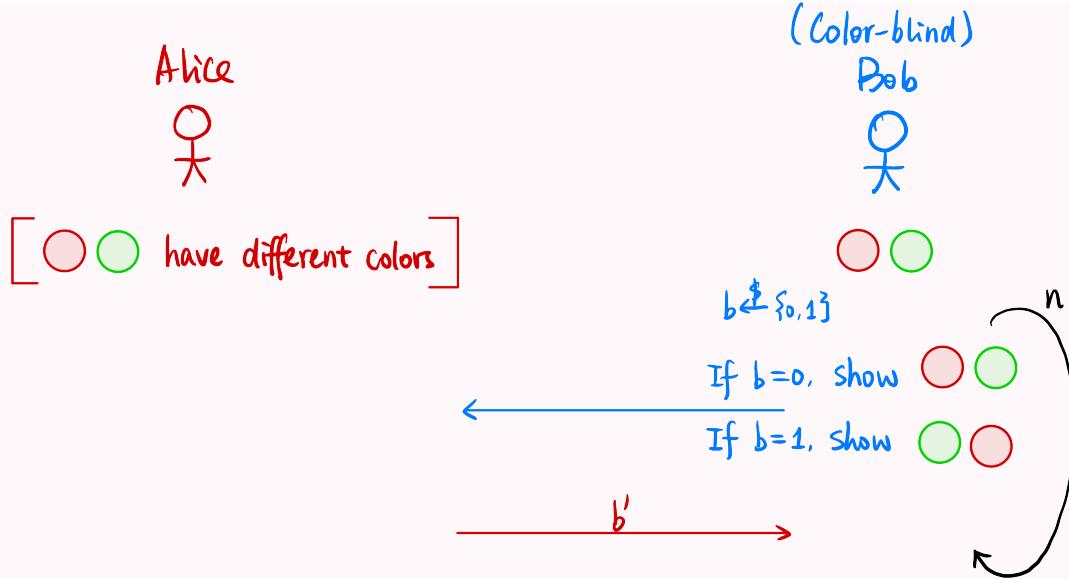
- Prove she knows the difference in taste between Coke and Pepsi without revealing how
- Prove that you have a bug in your code without revealing the bug
- She has the secret key for this ciphertext without revealing the plaintext
- Prove that she owns two different colored pens to her colorblind friend Bob

How is this possible? Let's examine this first scenario, which closely follows the red-green pens example from class *orthered – greenballsproblem*.

Example

Alice claims to have two different colored markers. She wants to prove this to her friend Bob, but there is an issue - Bob is colorblind!

Bob will randomly sample a bit $b \xleftarrow{\$} \{0, 1\}$, with $b = 0$ being "Stay" and $b = 1$ being "Swap". Bob will then either swap the order of the markers and show them to Alice, or keep them the same. Alice will give a guess b' of whether Bob swapped or not.



If statement is true: $\Pr[b=b']=1$

If statement is false: $\Pr[b=b']=\left(\frac{1}{2}\right)^n$

If the statement is true, $\Pr[b'=b]=1$ (Alice always gives the correct prediction).

If the statement is false, $\Pr[b'=b]=\frac{1}{2}$ (Alice is guessing with 0.5 probability).

To enhance this, we can run this a total of k times. If we run it enough times, Bob will be more and more confident in believing this. Alice getting this correct by chance has a $\frac{1}{2^k}$ probability.

The key idea, however, is that Bob doesn't gain any knowledge of how Alice differentiates.

Remark. This is a similar strategy in proving graph non-isomorphism.

For people who have seen this before, generally speaking, any NP language can be proved in zero-knowledge. Alice has the *witness* to the membership in NP language.

§1.5 Secure Multi-Party Computation

Alice



$$x \in \{0, 1\}$$

Second date? $y \in \{0, 1\}$

$$f(x, y) = x \wedge y$$

$$x \in \{0, 1\}^{1000}$$

Who is richer? $y \in \{0, 1\}^{1000}$

$$f(x, y) = \begin{cases} \text{Alice if } x > y \\ \text{Bob otherwise} \end{cases}$$

Bob



$$X = \left\{ \begin{array}{l} \text{friend}_A^1 \\ \vdots \\ \text{friend}_A^n \end{array} \right\}$$

Common friends?

$$f(X, Y) = X \wedge Y$$

$$Y = \left\{ \begin{array}{l} \text{friend}_B^1 \\ \vdots \\ \text{friend}_B^m \end{array} \right\}$$

$$X = \left\{ \begin{array}{l} (\text{username}, \text{password}) \\ \vdots \end{array} \right\}$$

$$Y = \left\{ \begin{array}{l} (\text{usr}, \text{psw}) \\ \vdots \end{array} \right\}$$

Example (Secure AND)

Alice and Bob go on a first date, and they want to figure out whether they want to go on a second date. They will only go on a second date if and only if both agree to a second date.

How will they agree on this? They could tell each other, but this could be embarrassing. One way is for them to share with a third-party (this is what dating apps do!). However, there might not always be an appropriate third party (in healthcare examples, not everyone can be trusted with the data).

In this case, Alice has a choice bit $x \in \{0, 1\}$ and Bob has a choice bit $y \in \{0, 1\}$. They are trying to jointly compute $f(x, y) = x \wedge y$.

Remark 1.4. Couldn't a party still figure out how the other party feels? For example, if Bob's bit was 1 and the joint result was 0, Bob can *infer* that Alice's bit was 0.

This is, in effect, the best we can do. The ideal guarantee is that each party only learns any information they can infer from the *output* and their input. However, they should not learn anything more.

What are we trying to achieve here? We want to jointly compute some function, where each party has private input, such that each party only learns the output. They should not learn anything about other parties' inputs.

Example (Yao's Millionaires' Problem)

Perhaps, Alice and Bob want to figure out who is richer. The inputs are $x \in \{0, 1\}^{1000}$ and $y \in \{0, 1\}^{1000}$ (for simplification, let's say they can express their wealth in 1000 bits). The output is the person who has the max.

$$f(x, y) = \begin{cases} \text{Alice} & \text{if } x > y \\ \text{Bob} & \text{otherwise} \end{cases}$$

Example (Private Set Intersection)

Alice and Bob meet for the first time and want to determine which of their friends they share. However, they do not want to reveal who specifically are their friends.

X is a set of A's friends $X = \{\text{friend}_A^1, \text{friend}_A^2, \dots, \text{friend}_A^n\}$ and Bob also has a set $Y = \{\text{friend}_B^1, \text{friend}_B^2, \dots, \text{friend}_B^m\}$. They want to jointly compute

$$f(X, Y) = X \cap Y.$$

You might need to reveal the cardinality of these sets, but you could also pad them up to a maximum number of friends.

This has a lot of applications in practice! In Google Chrome, your browser will notify you that your password has been leaked on the internet, without having access to your passwords in the clear. X will be a set of *your* passwords, and Google will have a set Y of *leaked* passwords. The *intersection* of these sets are which passwords have been leaked over the internet, without revealing all passwords in the clear.

Question. Isn't the assumption that the size of input records is revealed weaker than using a trusted third-party?

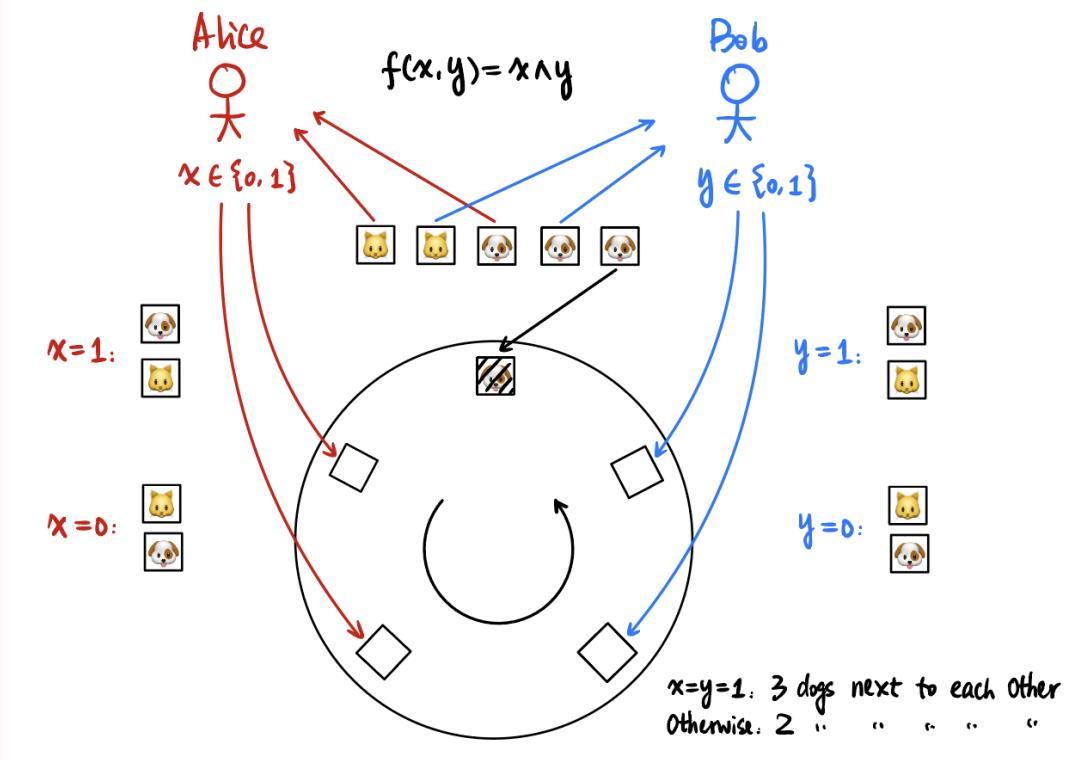
Yes, however in some cases (hospital health records), parties are legally obliged to keep data secure. We wish for security more than the secrecy of cardinality.

In the general case, Alice and Bob have some inputs x and y with bounded length, and they want to jointly compute some function f on these inputs. This is Secure Two-Party Computation. Furthermore, there could be multiple parties x_1, \dots, x_n that jointly compute $f(x_1, \dots, x_n)$ that hides each input. This is Secure Multiparty Computation.

We'll explore a toy example with the bit-AND from the dating example.

Example (Private Dating)

Alice and Bob have choice bits $x \in \{0, 1\}$ and $y \in \{0, 1\}$ respectively. There is a *physical* round table with 5 identical slots, one already filled in with a *dog* facing down. Alice and Bob each have identical *dog*, *cat* cards (each of the *dog* and *cat* cards are indistinguishable from cards of the same value). Note that our use of *dog*, *cat* is arbitrary, and could be any other set of two elements: 0, 1, x, o , etc.



Alice places her cards on the 2 slots in some order, and Bob does the same.

They then spin the table around and reveal all the cards, learning $x \wedge y$.

If $x = 1$, Alice places it as *dog* on top of *cat*, and if $y = 1$, Bob places it as *dog* on top of *cat* as well. Otherwise, they flip them. If $x = y = 1$, then the *dog*'s will be adjacent. If $x \neq y$, the order will be *dog, dog, cat, dog, cat* (the *cat*'s are not adjacent), regardless of which of Alice or Bob produced $x = 0$ (or both!).

This is a toy example! It doesn't use cryptography at all! Two parties have to sit in front of a table. This is called card-based cryptography. We will be using more secure primitives.

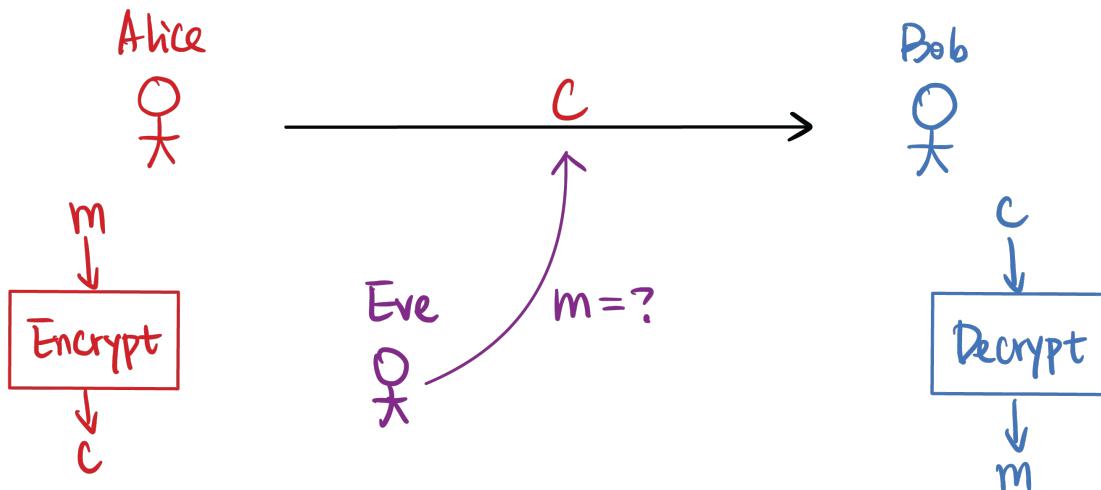
§1.6 Fully Homomorphic Encryption

We'll come back to the secure messaging example.

Alice wants to send Bob a message. She encrypts it somehow and sends a ciphertext $c_1 = \text{Enc}(m_1)$. A nice feature for some encryption schemes is for Eve to do some computation homomorphically on the ciphertexts. Eve might possibly want to add ciphertexts (that leads to plaintext adding)

$$c_1 = \text{Enc}(m_1), c_2 = \text{Enc}(m_2) \Rightarrow c' = \text{Enc}(m_1 + m_2)$$

or perhaps $c'' = \text{Enc}(m_1 \cdot m_2)$, or compute arbitrary functions. *Sometimes*, this is simply adding $c_1 + c_2$, but usually not.

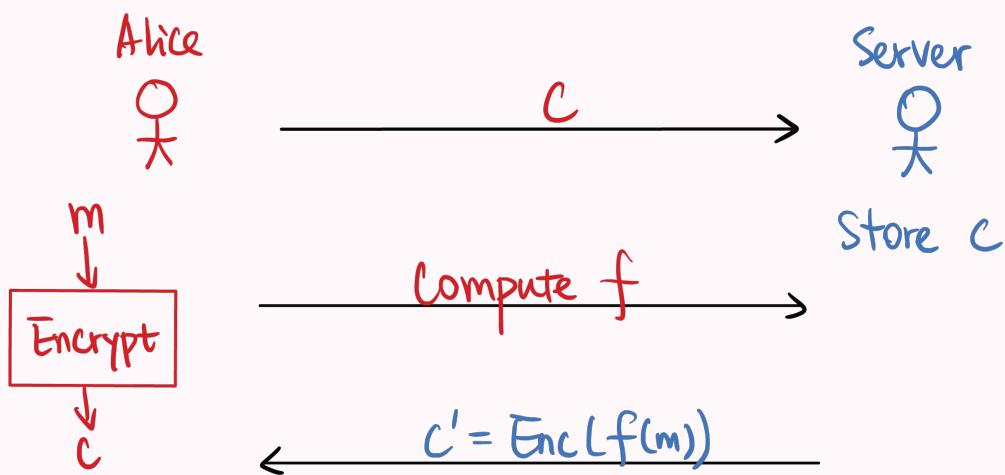


$$\begin{aligned} c_1 &= \text{Enc}(m_1) \\ c_2 &= \text{Enc}(m_2) \end{aligned} \Rightarrow \begin{aligned} c' &= \text{Enc}(m_1 + m_2) \\ c'' &= \text{Enc}(m_1 \cdot m_2) \end{aligned}$$

We want to hopefully compute any function in polynomial time!

Example (Outsourced Computation)

Alice has some messages but doesn't have enough compute. There is a server that has *a lot* of compute!



Alice encrypts her data and stores it in the server. At some point, Alice might want to compute a function on the encrypted data on the server, without the server revealing the original data. For example, she may want to search across that data without the server decrypting it (or sending her the entire encrypted DB). With homomorphic encryption, the server may be able to query and return an encrypted value to Alice without learning anything about the query itself.

This is an example of how fully homomorphic encryption can be useful.

Remark. This problem was not solved until 2009 (when Peihan started her undergrad). Theoretically, it doesn't even seem that possible! Being able to compute functions on ciphertexts that correspond to functions on plaintexts.

To construct fully-homomorphic encryption, we'll be using lattice-based cryptography which is a post-quantum secure!

§1.7 Further Topics

We might cover some other topics:

- Differential Privacy
- Crypto applications in machine learning
- Crypto techniques used in the blockchain²

²One important technique is Zero-Knowledge proofs, for example.

- LLM Watermarking

What else would you like to learn? What else do you want to understand? Do go through the semester with these in mind! How do I log into Google? How do I send messages to friends?

Feel free to let us know on Ed!

§1.8 Q & A

- *Do I need to have a crypto background?*

No!

- *Why C++*

Existing crypto libraries are mostly in C++ and most students have seen C/C++ in either cs33 or cs300. We did, however, consider implementing everything in Rust!

- *Class Participation*

Course attendance is expected for all students, since a large portion of the class structure relies on students answering questions in class. Even if you don't actively volunteer, try to think through the answer on your own! Students may join class via Zoom if they are remote-only, traveling, or sick.

- *What is the difference between CSCI 1515 and CSCI 1510, MATH 1580, or CSCI 1040?*

CSCI 1510 is essentially “theoretical cryptography.” It covers formal definitions and constructions and proofs. There is no coding, just proofs.

MATH 1580 considers crypto from the mathematical perspective. They try to understand some of the computational assumptions we assume from a mathematical standpoint. I.e. why is factoring hard to compute, and what is the best algorithm to compute it? In CS, we simply assume factoring is hard. MATH 1580 is more similar to number theory and group theory.

CSCI 1040 is a cryptography course for students without a computer science or math background. It is self described as "crypto for poets!"

CSCI 1515, on the other hand, takes a hands on approach to cryptography. This class is suitable for both students with limited cryptography exposure more interested in cryptography from a software engineering perspective, as well as students with a cryptography theory background looking to understand what makes an implementation efficient in practice.

This course **will not** be offered in Spring 2027.

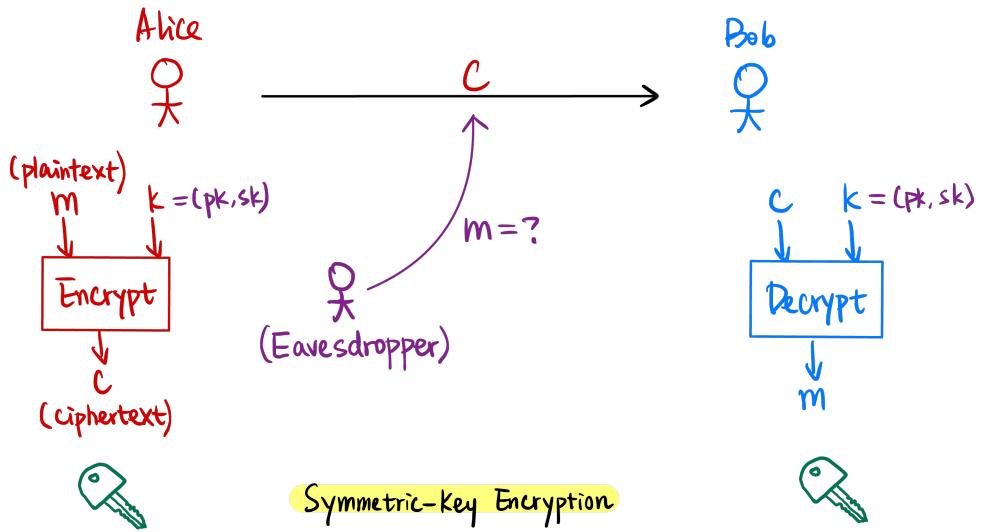
§2 January 28, 2026

Make-up lecture for our canceled Monday class will take place on Friday 3-4:20 pm in Friedman Hall 202 and on Zoom.

§2.1 Encryption Scheme Basics

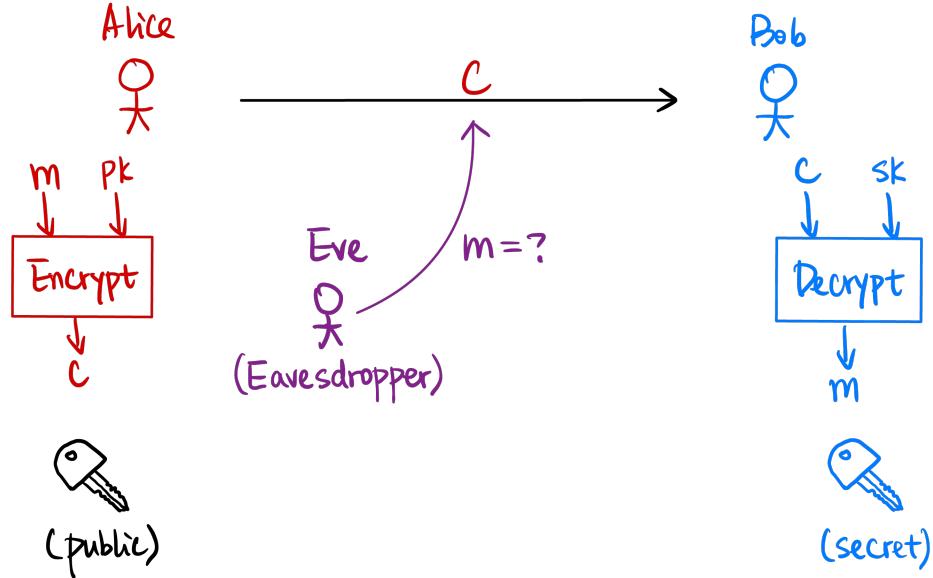
This lecture we'll cover encryption schemes. We briefly mentioned what encryption schemes were last class, we'll dive into the technical content: how we construct them, assumptions, RSA, ElGamal, etc.

Fundamentally, an encryption scheme protects message secrecy. If Alice wants to communicate to Bob, Alice will encrypt a message (plaintext) using some key which gives her a ciphertext. Sending the ciphertext through Bob using a public channel, Bob can use the key to decrypt the ciphertext and recover the message. An eavesdropper in the middle will have no idea what message has been transmitted.



In this case, they are using a shared key, which we called secret-key encryption or symmetric-key encryption.

A stronger version of private-key encryption is called public-key encryption. Alice and Bob do not need to agree on a shared secret key beforehand. There is a keypair (pk, sk) , a *public* and *private* key.



§2.1.1 Syntax

Definition 2.1 (Symmetric-Key Encryption)

A symmetric-key encryption (SKE) scheme contains 3 algorithms, $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$.

Generation. Generates key $k \leftarrow \text{Gen}$.

Encryption. Encrypts message m with key k , $c \leftarrow \text{Enc}(k, m)$, which we sometimes write as $\text{Enc}_k(m)$.

Decryption. Decrypts using key k to retrieve message m , $m := \text{Dec}(k, c)$, or written as $\text{Dec}_k(c)$.

Note the notation \leftarrow and $:=$ is different. In the case of generation and encryption, the produced key k or c follows a *distribution* (is randomly sampled), while we had better want decryption to be deterministic in producing the message.

In other words, we use \leftarrow when the algorithm might involve randomness and $:=$ when the algorithm is deterministic.

Definition 2.2 (Public-Key Encryption)

A public-key encryption (PKE) scheme $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$ contains the same 3 algorithms,

Generation. Generate keys $(pk, sk) \leftarrow \text{Gen}$.

Encryption. Use the public key to encrypt, $c \leftarrow \text{Enc}(pk, m)$ or $\text{Enc}_{pk}(m)$.

Decryption. Use the secret key to decrypt, $m := \text{Dec}(pk, c)$ or $\text{Dec}_{sk}(c)$.

Remark 2.3. Note that all these algorithms are public knowledge. This is known as Kerckhoff's principle.

Intuitively, one key reason is because if the security of our scheme relies on hiding the *algorithm*, then if it is leaked we need to construct an entirely new algorithm. However, if the security of our scheme relies on, for example, a secret key, then we simply need to generate a new key.

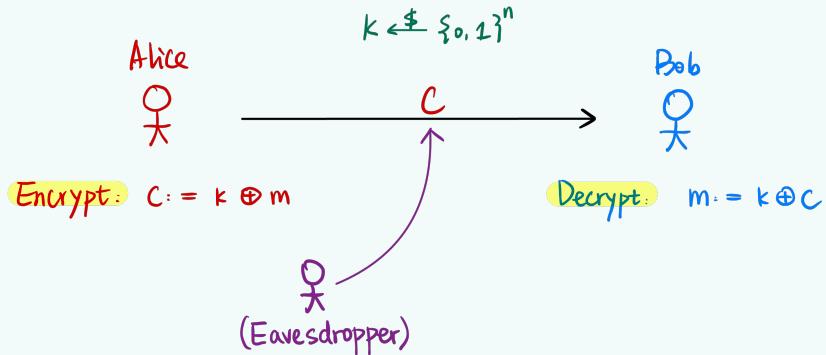
Question. If we can construct public-key encryption, why do we even bother with secret-key encryption? We could just use the (pk, sk) pair for our secret key, and this does the same thing.

1. First of all, public-key encryption is almost always *more expensive*. Symmetric-key encryption schemes give us efficiency.
2. Public-key encryption relies on much stronger computational assumptions, whereas symmetric key encryption are still post-quantum secure.

§2.1.2 Symmetric-Key Encryption Schemes

Definition 2.4 (One-Time Pad)

Secret key is a uniformly randomly sampled n bit string $k \xleftarrow{\$} \{0, 1\}^n$.



Encryption. Alice uses the secret key and bitwise-XOR with the plaintext.

$$\begin{array}{rcl} \text{secret key } & k = 0100101 \\ \oplus \text{ plaintext } & m = 1001001 \\ \hline \text{ciphertext } & c = 1101100 \end{array}$$

Decryption. Bob uses the secret key and again bitwise-XOR with the ciphertext

$$\begin{array}{rcl} \text{secret key } & k = 0100101 \\ \oplus \text{ ciphertext } & c = 1101100 \\ \hline \text{plaintext } & m = 1001001 \end{array}$$

This is widely used in cryptography, called *masking* or *unmasking*.

Question. Why is this correct?

An XOR done twice with the same choice bit b is the identity. In other words, $k \oplus (k \oplus m) = m$

Question. Why is this secure?

We can think about this as the distribution of c . $\forall m \in \{0, 1\}^n$, the encryption of m is uniform over $\{0, 1\}^n$ (since k was uniform).

Another way to think about this is that for any two messages $m_0, m_1 \in \{0, 1\}^n$, $\text{Enc}_k(m_0) \equiv \text{Enc}_k(m_1)$. That is, the encryptions follow the *exact same* distribution. In this case, they are both uniform, but this is not always the case.

Question. Can we reuse k ? Should we use the same key again to encrypt another message? Or, it is possible for the eavesdropper to extract information.

For example, $\text{Enc}_k(m)$ is $c := k \oplus m$, and $\text{Enc}_k(m')$ is $c' := k \oplus m'$. If the two messages are the same, the ciphertexts are the same.

By XOR c and c' , we get

$$\begin{aligned} c \oplus c' &= (k \oplus m) \oplus (k \oplus m') \\ &= m \oplus m' \end{aligned}$$

This is why this is an *one-time pad*. This is a bit of an issue, to send an n -bit message, we need to agree on an n -bit message.

In fact, this is *the best* that we can do.

Theorem 2.5 (Shannon's Theorem)

Informally, for perfect (information-theoretic³) security, the key space must be at least as large as the message space.

$$|\mathcal{K}| \geq |\mathcal{M}|$$

where \mathcal{K} is the key space and \mathcal{M} is the message space.

Question. How can we circumvent this issue?

The high level idea is that we weaken our security guarantees *a little*. Instead of saying that they have to be *exactly the same* distribution, we say that they are *hard to distinguish* for an adversary with limited computational power. This is how modern cryptography gets around these lower bounds in classical cryptography. We can make *computational assumptions* about cryptography.

We can think about computational security,

Definition 2.6 (Computational Security)

We have computational security when two ciphertexts have distribution that cannot be distinguished using a polynomial-time algorithm.

Definition 2.7 (Polynomial-Time Algorithm)

A polynomial time algorithm $A(x)$ is one that takes input x of length n , A 's running time is $O(n^c)$ for a constant c .

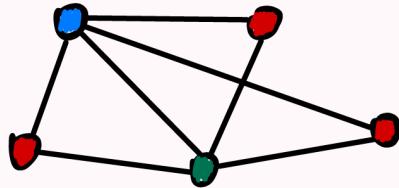
³That the distributions of ciphertexts are identical, that $\text{Enc}_k(m_0) \equiv \text{Enc}_k(m_1)$.

Definition 2.8 (NP Problem)

A decision problem is in nondeterministic polynomial-time when its solution can be *verified* in polynomial time.

Example 2.9 (Graph 3-Coloring)

Given a graph, does it have a 3-coloring such that no two edges join the same color? For example,



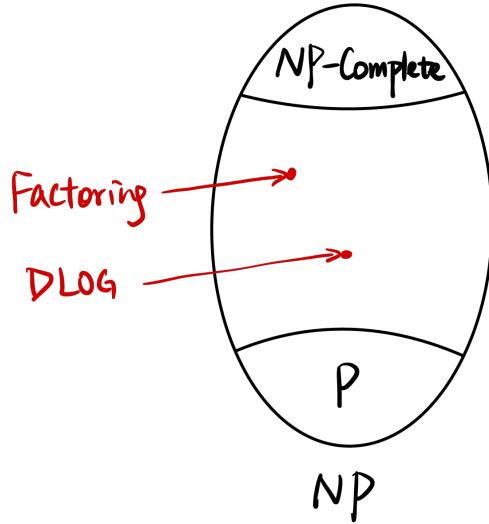
This can be *verified* in polynomial time (we can check if such a coloring is a valid 3-coloring), but it is computed in NP time.

Definition 2.10

An NP-complete problem is a “hardest” problem in NP. Every problem in NP is at least as hard as an NP-complete problem.

Right now, we assume $P \neq NP$. As of right now, there is no realistic algorithm that can solve any NP problem in polynomial-time.

Even further, we pick some problems not in NP-complete, not in P. We assume they are neither NP-complete nor in P (we don’t yet have a reduction, but we don’t know if one could exist) The reasoning behind using these problems is as we have no good cryptoscheme relying solely on NP-complete problems (we need something weaker).



Going back to our definition of computational security [definition 2.6](#),

Definition (Computational Security)

Let the adversary be computationally bounded (i.e. only a *polynomial-time algorithm*). Then \forall probabilistic poly-time algorithm \mathcal{A} ,

$$\text{Enc}_k(m_0) \stackrel{c}{\sim} \text{Enc}_k(m_1)$$

Where $\stackrel{c}{\sim}$ is “computationally indistinguishable”.

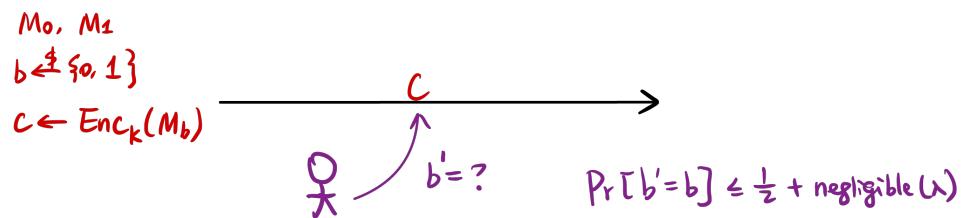
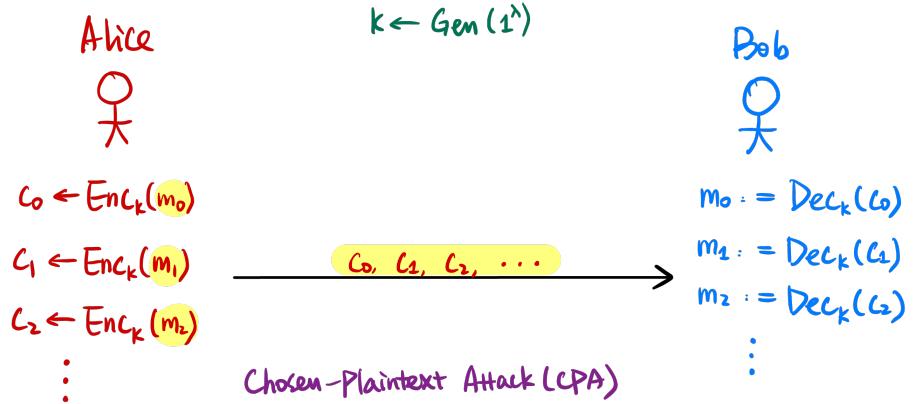
What does it mean for distributions to be “computationally indistinguishable”? Let’s say Alice encrypts multiple messages m_0, m_1, \dots to Bob and produces c_0, c_1, \dots . Even if Eve can see all plaintexts m_i in the open and ciphertexts c_i in the open, between known m_0, m_1 and randomly encrypting one of them $c \leftarrow \text{Enc}_k(m_b)$ where $b \xleftarrow{\$} \{0, 1\}$, the adversary cannot determine what the random choice bit b is. That is, $\Pr[b = b'] \leq \frac{1}{2} + \text{negligible}(\lambda)$ ⁴. This is Chosen-Plaintext Attack (CPA) Security.

Definition 2.11 (Chosen-Plaintext Attack)

In simpler terms, a Chosen-Plaintext Attack (CPA) allows an adversary to request the encryptions of any number of chosen messages with the same key. The adversary can see the ciphertext of each message. Then, the adversary chooses two messages to be encrypted, and receives the ciphertext of one. The adversary must try to determine which message the ciphertext corresponds to.

Note: the adversary is allowed to send the same message as many times as it wishes.

⁴ λ is the security parameter, roughly a measure of how secure the protocol is. If it were exactly equal $\frac{1}{2}$, we have information-theoretic security.



For a key generated $k \leftarrow \text{Gen}(1^\lambda)$.

Theoretically, for λ a security parameter and an adversary running in time $\text{poly}(\lambda)$, the adversary should have distinguishing advantage $\text{negligible}(\lambda)$ where

$$\text{negligible}(\lambda) \ll \frac{1}{\lambda^c} \quad \forall \text{ constant } c.$$

In practice, we set $\lambda = 128$. This means that the best algorithm to break the scheme (e.g. find the secret key) takes time $\sim 2^\lambda$. Currently, this is longer than the age of the universe.

Remark 2.12. Just how big is 2^{128} ? Well, see how long 2^{128} CPU cycles will take.

Let's assume the CPU spec is 3.8 GHz i.e. $3 \cdot 10^9$ cycles per second. Moreover, note that $2^{128} \sim 10^{40}$ CPU cycles.

Then doing the math ...

$$10^{40} \text{ CPU cycles} \cdot \frac{1s}{10^9 \text{ CPU cycles}} \cdot \frac{1 \text{ year}}{31,536,000s} \sim 10^{22} \text{ years}$$

Now let's be generous and say the age of the universe is 10^{11} years ... then we see that it would still take 10^{11} times the current age of the universe!

Example 2.13

If the best algorithm is a brute-force search for k , what should our key length be?

It can just be a λ bit string.

Example

What if the best algorithm is no longer a brute-force search, but instead for a key length l takes $\sim \sqrt{2^l}$?

Our key length should be 2λ . Doing the math, we want $\sqrt{2^l} \equiv 2^\lambda$, solving for l gives 2λ .

Going back to the original problem of secret-key encryption, how can we use our newfound cryptographic constructions to improve this?

From a pseudorandom function/permuation (PRF/PRP), we can reuse our secret key by passing it through the pseudorandom function.

The current practical construction for PRD/PRP is called the block cipher, and the standardized implementation is AES⁵

It is a computational assumption⁶ that the AES construction is secure, and the best attack is currently a brute-force search (in both classical and quantum computing realms).

§2.1.3 Public-Key Encryption Schemes

Using computational assumptions, we explore some public-key encryption schemes.

RSA Encryption. This is based on factoring/RSA assumption, that factoring large numbers is hard.

EIGamal Encryption. This is based on the discrete logarithm/Diffie-Hellman Assumption, that finding discrete logs in \mathbb{Z}_p is hard.

Lattice-Based Encryption. The previous two schemes are not quantum secure. Quantum computation will break these schemes. Lattice-based encryption schemes are post-quantum secure. They are associated with the difficulty of finding ‘short’ vectors in lattices⁷.

⁵Determined via a competition for such an algorithm in the early 2000s.

⁶Based on heuristic, not involving any number theory!

⁷Covered later in class, we focus on the first two now.

Theorem 2.14

(Very informally,) It is impossible to construct PKE from SKE in a black-box way. This is called “black-box separation”.

We first need to define a bit of number theory background.

Definition 2.15

We denote $a \mid b$ as a divides b , that is, there is integer c such that $b = a \cdot c$.

Definition 2.16 (Primes)

An integer $p > 1$ that only has two divisors: 1 and p

Definition 2.17 (Mod)

$a \bmod N$ is the remainder of a when divided by N .

$a \equiv b \pmod{N}$ means when a and b are congruent modulo N . That is, $a \bmod N = b \bmod N$.

Question. How might we compute $a^b \bmod N$? What is the time complexity? Let a, b, N be n -bit strings.

Naïvely, we can repeatedly multiply. But this takes b steps (2^n).

We can ‘repeatedly square’. For example, we can get to a^8 faster by getting a^2 , squaring to get a^4 , and again to get a^8 . We can take the bitstring of b and determine how to compute this.

Example

If $b = 100101_2$, we take $a \cdot a^4 \cdot a^{32} \bmod N$ which can be calculated recursively (an example is given in the first assignment).

The time complexity of this is order $O(n)$ for n -bit a, b, N ⁸.

Definition 2.18

The $\gcd(a, b)$ is the greatest common divisor of a, b . If $\gcd(a, b) = 1$, then a, b are coprime.

⁸Not exactly order n , we should add the complexity of multiplication. However, this should be bounded by N since we can log at every step.

Question. How do we compute gcd? What is its time complexity?

Example

We use the Euclidean Algorithm. Take $\gcd(12, 17)$,

$$\begin{aligned} 17 \mod 12 &= 5 \\ 12 \mod 5 &= 2 \\ 5 \mod 2 &= 1 \\ 2 \mod 1 &= 0 \end{aligned}$$

or take $\gcd(12, 18)$

$$\begin{aligned} 18 \mod 12 &= 6 \\ 12 \mod 6 &= 0 \end{aligned}$$

Theorem 2.19 (Bezout's Theorem, roughly)

If $\gcd(a, N) = 1$, then $\exists b$ such that

$$a \cdot b \equiv 1 \pmod{N}.$$

This is to say, a is invertible modulo N . b is its inverse, denoted as a^{-1} .

Question. How do we compute b ?

We can use the Extended Euclidean Algorithm!

Example

We write linear equations of a and N that sum to 1, using our previous Euclidean Algorithm. Take the previous example $\gcd(12, 17)$,

$$\begin{aligned} 17 \mod 12 &= 5 \\ 12 \mod 5 &= 2 \\ 5 \mod 2 &= 1 \\ 2 \mod 1 &= 0 \end{aligned}$$

We write this as

$$\begin{aligned} 5 &= 17 - 12 \cdot 1 \\ 2 &= 12 - 5 \cdot 2 = 12 \cdot x + 17 \cdot y \\ 1 &= 5 - 2 \cdot 2 = 12 \cdot x' + 17 \cdot y' \end{aligned}$$

where we substitute the linear combination of 5 into 5 on line 2, substitute linear combination of 2 into 2 on line 1, each producing another linear combination of 12, 17.

If $\gcd(a, N) = 1$, we use the Extended Euclidean Algorithm to write $1 = a \cdot x + N \cdot y$, then $1 \equiv a \cdot x \pmod{N}$.

Definition 2.20 (Group of Units mod N)

We have set

$$\mathbb{Z}_N^\times := \{a \mid a \in [1, N-1], \gcd(a, N) = 1\}$$

which is the group of units modulo N (they are units since they all have an inverse by above).

Definition 2.21 (Euler's Phi Function)

Euler's phi (or totient) function, $\phi(N)$, counts the number of elements in this set. That is, $\phi(N) = |\mathbb{Z}_N^\times|$.

Theorem 2.22 (Euler's Theorem)

For all a, N where $\gcd(a, N) = 1$, we have that

$$a^{\phi(N)} \equiv 1 \pmod{N}.$$

With this, we can start talking about RSA.

§2.1.4 RSA

We first define the RSA assumption.

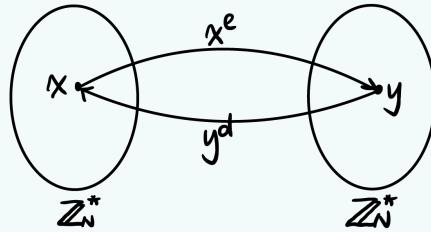
Definition 2.23 (Factoring Assumption)

Given two n -bit primes p, q , we compute $N = p \cdot q$. Given N , it's computationally hard to find p and q (classically).

Definition 2.24 (RSA Assumption)

Given two n -bit primes, we again compute $N = p \cdot q$, where $\phi(N) = (p - 1)(q - 1)$. We choose an e such that $\gcd(e, \phi(N)) = 1$ and compute $d \equiv e^{-1} \pmod{\phi(N)}$.

Given N and a random $y \xleftarrow{\$} \mathbb{Z}_N^*$, it's computationally hard to find x such that $x^e \equiv y \pmod{N}$.



However, given p, q , it's easy to find d . We know $\phi(N) = (p - 1)(q - 1)$, so we can compute d from e by running the Extended Euclidean Algorithm. Then, taking $(x^e)^d \equiv x^{ed} \equiv x$ which allows us to extract x again.

Encrypting is exactly raising by power d , and decrypting is raising again by power e .

Remaining questions:

- How can we generate primes p, q ?
- How can we pick e such that $\gcd(e, \phi(N)) = 1$?
- What security issues can you see?

We'll continue next class.

§3 January 30, 2026

§3.1 Basic Number Theory, *continued*

We recall a couple more definitions.

Definition 3.1

We first define the multiplicative group of integers modulo n as

$$\mathbb{Z}_N^* = \{a \in [1, N - 1] \mid \gcd(a, N) = 1\}$$

Definition 3.2

We define the Euler's totient function as

$$\phi(N) = |\mathbb{Z}_N^*|$$

Example 3.3

If $N = p \cdot q$ where p, q are prime, then $\phi(N) = (p - 1)(q - 1)$.

Theorem 3.4 (Euler's Theorem)

$\forall a, N$ where $\gcd(a, N) = 1$, we have that $a^{\phi(N)} \equiv 1 \pmod{N}$.

Corollary 3.5

If

$$d \equiv e^{-1} \pmod{\phi(N)}$$

, then

$$\forall a \in \mathbb{Z}_N^*, (a^d)^e \equiv a \pmod{N}$$

.

§3.2 RSA Encryption, *continued*

Recall: that the RSA encryption algorithm contains 3 components:

Gen(1^λ): Generate two n -bit primes p, q . We compute $N = p \cdot q$ and $\phi(N) = (p - 1)(q - 1)$. Choose e such that $\gcd(e, \phi(N)) = 1$. We compute $d = e^{-1} \pmod{\phi(N)}$. Our public key $pk = (N, e)$, our secret key is $sk = d$.

$\text{Enc}_{pk}(m)$: $c = m^e \pmod{N}$.

$\text{Dec}_{sk}(c)$: $m = c^d \pmod{N}$.

We have a few remaining questions:

1. How do we generate 2 primes p, q ? More specifically, how do we generate two *large* primes efficiently?
2. How do we choose such an e ?
3. How do we compute $d = e^{-1} \pmod{\phi(N)}$?
4. How do we efficiently compute $m^e \pmod{N}$ and $c^d \pmod{N}$?

Question. How do we resolve these issues to ensure the `Gen` step is efficient (polynomial time)?

1. We pick an arbitrary number p and check for primality efficiently (using Miller Rabin, a probabilistic primality test). We pick random numbers until they are prime. Since primes are ‘pretty dense’ in the integers, this can be done efficiently.
2. Since we’re unsure whether coprime numbers are dense, we can just pick a small prime e . Guessing, like we did with p, q is also valid, although not strictly necessary.
3. We can compute d using the Extended Euclidean Algorithm.
4. We can repeatedly square (using fast power algorithm).

Question. What happens if we can factor N ?

Then we can find p and q and calculate $\phi(N) = (p - 1)(q - 1)$, and then we can compute $d = e^{-1} \pmod{\phi(N)}$. Thus, RSA relies crucially on the factoring problem being hard.

Question. The above scheme is known as “plain” RSA. Are there any security issues?

- It relies on factoring being difficult (this is the computational assumption). Post-quantum, Shor’s Algorithm will break RSA.
- Recall last lecture that CPA (Chosen-Plaintext Attack) security was defined as an adversary not being able to discern between an encryption of m_0 and m_1 , *knowing* m_0 and m_1 in the clear.

Eve could just encrypt m_0 and m_1 themselves using public e , and discern which of the plaintexts the ciphertext corresponds to. For RSA, this is a very concrete attack.

The concrete reason is that the encryption algorithm Enc is *deterministic*. If you encrypt the same message twice, it will be the same ciphertext. Since this scheme is deterministic, it fails CPA security - for example, an adversary can encrypt messages and look for matches. We really want to be sure that $m \xleftarrow{\$} \mathbb{Z}_N^\times$ (that it has enough entropy).

Question. In practice, how can RSA be useful with these limitations?

As long as we pick the plaintext which is randomly sampled, security for RSA holds. There is also a more involved way of using RSA that is CPA-secure, but we will not go in detail of it.

Remark. In practice, we usually set length of p and q to be 1024 bits, and the key length is 2048 bits.

Moreover, note that although exponentiation can be done in polynomial time, it's still a very expensive operation. This is why public key encryption is, in general, more expensive than symmetric key encryption.

§3.3 Intro to Group Theory

Definition 3.6 (Group)

A group is a set \mathbb{G} along with a binary operation \circ with properties:

Closure. $\forall g, h \in \mathbb{G}, g \circ h \in \mathbb{G}$.

Existence of an identity. $\exists e \in \mathbb{G}$ such that $\forall g \in \mathbb{G}, e \circ g = g \circ e = g$.

Existence of inverse. $\forall g \in \mathbb{G}, \exists h \in \mathbb{G}$ such that $g \circ h = h \circ g = e$. We denote the inverse of g as g^{-1} .

Associativity. $\forall g_1, g_2, g_3 \in \mathbb{G}, (g_1 \circ g_2) \circ g_3 = g_1 \circ (g_2 \circ g_3)$.

We say a group is additionally *Abelian* if it satisfies commutativity

Commutativity. $\forall g, h \in \mathbb{G}, g \circ h = h \circ g$.

For a finite group, we use $|\mathbb{G}|$ to denote its *order*.

Example 3.7

$(\mathbb{Z}, +)$ is an Abelian group.

We can check so: two integers sum to an integer, identity is 0, the inverse of a is $-a$, addition is associative and commutative.

(\mathbb{Z}, \cdot) is not a group. There is no inverse for 0 such the $0 \cdot h = 1$.

$(\mathbb{Z}_N^\times, \cdot)$ is an Abelian group (\cdot is multiplication mod N).

Definition 3.8 (Cyclic Group)

Let \mathbb{G} be a group of order m . We denote

$$\langle g \rangle := \{e = g^0, g^1, g^2, \dots, g^{m-1}\}.$$

\mathbb{G} is a cyclic group if $\exists g \in \mathbb{G}$ such that $\langle g \rangle = \mathbb{G}$. g is called a generator of \mathbb{G} .

Example

\mathbb{Z}_p^\times (for prime p) is a cyclic group of order $p - 1$ ⁹.

$$\mathbb{Z}_7^\times = \{3^0 = 1, 3^1, 3^2 = 2, 3^3 = 6, 3^4 = 5, 3^5 = 5\}.$$

Question. How do we find a generator?

For every element, we can continue taking powers until $g^\alpha = 1$ for some α . We hope that $\alpha = p - 1$ (the order of g is the order of the group), but we know at least $\alpha \mid p - 1$.

§3.4 Computational Assumptions

We have a few assumptions we make called the Diffie-Hellman Assumptions, in order of **weakest to strongest**¹⁰ assumptions.

Let $(\mathbb{G}, q, g) \leftarrow \mathcal{G}(1^\lambda)$ be a cyclic group \mathbb{G} of order q (a $\theta(\lambda)$ -bit integer) with generator g . For integer groups, keys are usually 2048-bits. For elliptic curve groups, keys are usually 256-bits.

Definition 3.9 (Discrete Logarithm (DLOG) Assumption)

Let $x \xleftarrow{\$} \mathbb{Z}_q$. We compute $h = g^x$.

Given (\mathbb{G}, q, g, h) , it's computationally hard to find the exponent x (classically).

⁹A proof of this extends beyond the scope of this course, but you are recommended to check out Math 1560 (Number Theory) or Math 1580 (Cryptography). You can take this on good faith.

¹⁰If one can solve DLOG, we can solve CDH. Given CDH, we can solve DDH. This is why CDH is *stronger* than DDH, and DDH is *stronger* than DLOG. It's not necessarily true the other way around (similar to factoring and DSA assumptions).

Definition 3.10 (Computational Diffie-Hellman (CDH) Assumption)

$x, y \xleftarrow{\$} \mathbb{Z}_q$, compute $h_1 = g^x, h_2 = g^y$.

Given $(\mathbb{G}, q, g, h_1, h_2)$, it's computationally hard to find g^{xy} .

Definition 3.11 (Decisional Diffie-Hellman (DDH) Assumption)

$x, y, z \xleftarrow{\$} \mathbb{Z}_q$. Compute $h_1 = g^x, h_2 = g^y$.

Given $(\mathbb{G}, q, g, h_1, h_2)$, it's computationally hard to distinguish between g^{xy} and g^z .

$$(g^x, g^y, g^{xy}) \stackrel{c}{\sim} (g^x, g^y, g^z).$$

§3.5 ElGamal Encryption

The ElGamal encryption scheme involves the following:

Gen(1^λ): We generate a group $(\mathbb{G}, q, g) \leftarrow \mathcal{G}(1^\lambda)$. We sample $x \xleftarrow{\$} \mathbb{Z}_q$, compute $h = g^x$. Our public key is $pk = (\mathbb{G}, q, g, h)$, secret key $sk = x$.

Enc _{pk} (m): We have $m \in \mathbb{G}$. We randomly sample $y \xleftarrow{\$} \mathbb{Z}_q$, which helps prevent our ciphertext from being deterministic. Our ciphertext is $c = \langle g^y, h^y \cdot m \rangle$. Note that $h = g^x$, so $g^{xy} \stackrel{c}{\sim} g^z$ is a one-time pad for our message m .

Dec _{sk} (c): To decrypt $c = \langle c_1, c_2 \rangle$, we raise

$$\begin{aligned} c_1^x &= (g^y)^x = g^{xy} \\ m &= \frac{g^{xy} \cdot m}{g^{xy}} = c_2 \cdot (c_1^x)^{-1}. \end{aligned}$$

Notes about ElGamal:

- Our group can be reused! We can use a public group that is fixed. In fact, there are *popular* groups out there used in practice. Some of these are Elliptic Curve groups which are much more efficient than integer groups. You don't need to use the details, yet you can use it! You can use any group, so long as the group satisfies the DDH assumption.
- Similar to RSA, this is breakable post-quantum. Given Shor's Algorithm, we can break discrete log.

§3.6 Secure Key Exchange

Using DDH, we can construct something very important, *secure key exchange*.

Definition 3.12 (Secure Key Exchange)

Alice and Bob sends messages back and forth, and at the end of the protocol, can agree on a shared key.

An eavesdropper looking at said communications cannot figure out what shared key they came up with.

Theorem 3.13

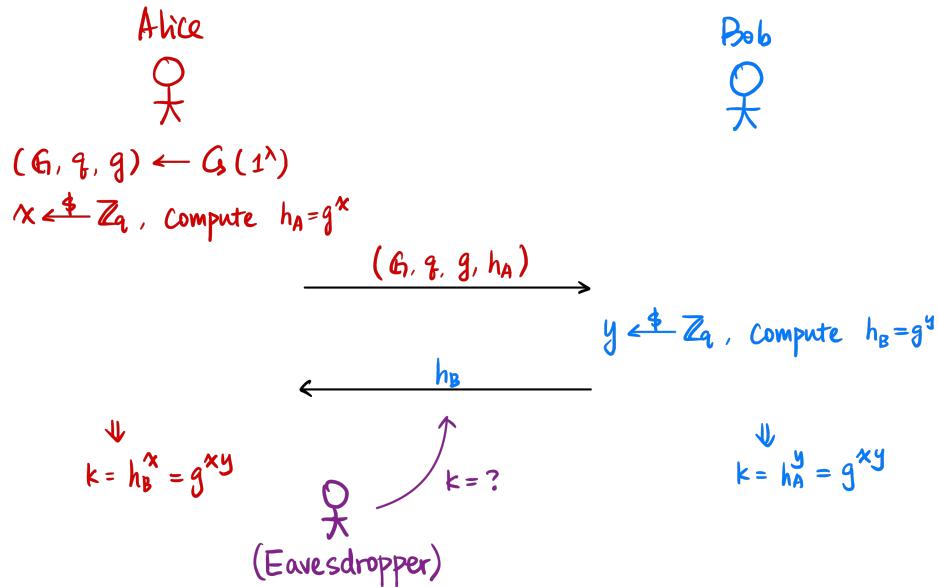
Informally, It's impossible to construct secure key exchange from secret-key encryption in a black-box way.

Question. How do we build a key exchange from public-key encryption?

Bob generates a keypair (pk, sk) . Alice generates a shared key $k \xleftarrow{\$} \{0, 1\}^\lambda$, and sends $\text{Enc}_{pk}(k)$ to Bob.

Using Diffie-Hellman, it's very easy. We have group $(\mathbb{G}, q, g) \leftarrow \mathcal{G}(1^\lambda)$. Alice samples $x \xleftarrow{\$} \mathbb{Z}_q$ and sends g^x . Bob also samples $y \xleftarrow{\$} \mathbb{Z}_q$ and sends g^y . Both Alice and Bob compute $g^{xy} = (g^x)^y = (g^y)^x$.

Diffie-Hellman Key Exchange



What happens in practice is that parties run Diffie-Hellman key exchange to agree on a shared key. Using that shared key, they run symmetric-key encryption. This gives us efficiency. Additionally, private-key encryptions don't rely on heavy assumptions on the security of protocols (such as the DDH, RSA assumptions).